

Introduction to Julia

Taylor Pospisil

2015-04-23

Motivation: The Two Language Problem

R/Python is too slow

- ▶ How do you get fast R code? Use Rcpp/Cython
- ▶ But it's easy to work with and you can work in the REPL
- ▶ Lots of libraries means someone has already done the work for you

C is too slow

- ▶ As fast as you're going to get.
- ▶ But don't forget about the programmer's time. It's not easy, and bugs are not easy to find.

Julia: A happy medium

- ▶ Best of both worlds: REPL with close-to-C speed

Concrete Example: Dynamic Programming

In the game of Risk, we want to calculate the probability of victory $p(A, D)$ attacking with A armies when the defender has D armies. Thus, we can use the following recursion

$$p(A, D) = \sum_{l=0}^A t(l, \min(A, 3), \min(D, 2)) \times p(A - l, D - (n - l))$$

Development

- ▶ Still a work-in-progress (current release = v0.37)
- ▶ Entire project on github.com/JuliaLang
- ▶ Package system (547 total) :
 - ▶ Gadfly (ggplot-like graphics)
 - ▶ PyCall, RCall
 - ▶ DataFrames
- ▶ Documentation available

Quick Technical Overview

JIT compilation

Code is compiled upon first execution

Type System

- ▶ The usual types included (ints, floats, arrays, dicts, sets, etc)
- ▶ User-defined types have the same performance as "built-ins"

Multiple Dispatch

Function calls are dispatched based upon the types of the arguments

Optimizing for Performance

Vectorization

Unlike with R, vectorization is **not** your first step

Type Stability

Very important, as type inference allows compilation of efficient code

The Usual

- ▶ Avoid memory allocation in tight loops
- ▶ Use the profiler
- ▶ Avoid pre-mature optimization and nothing makes bad algorithms fast

When is Julia > Matlab

Julia > Matlab

- ▶ Julia > R > Matlab
- ▶ Julia is open-source and free
- ▶ Julia has easy-to-read syntax
- ▶ Julia has good support for linear algebra

Matlab > Julia

- ▶ ??

Pick Julia when you don't have a license or don't rely upon existing Matlab code

When is Julia > Python

Julia > Python

- ▶ No python2/python3 issues
- ▶ No numpy issues; numerical computation isn't an add-on for Julia

Julia < Python

- ▶ Python has many preexisting libraries (beautifulsoup, sklearn, nltk)
- ▶ Python has a much bigger user/dev base

Pick Julia when you have to do something original and numerical

When is Julia > R

Julia > R

- ▶ Julia is better suited for general purpose tasks
- ▶ Julia is faster for iteration (MCMC)

R > Julia

- ▶ R has existing packages
- ▶ R has better graphics (for now)
- ▶ Everyone knows R

Julia > R when you want to do something original and computationally intensive (but without needing C)

Takeaway

- ▶ Julia shows promise from a performance standpoint
- ▶ Julia is open-source and free
- ▶ Packages aren't there (yet?), but might use RCall or PyCall
- ▶ If you need C-like performance but R-like ease of development, give Julia a try!